



R news and tutorials contributed by (750) R bloggers

- [Home](#)
- [About](#)
- [RSS](#)
- [add your blog!](#)
- [Learn R](#)
- [R jobs](#)
- [Contact us](#)

Welcome!

Follow @rbloggers { 53.4K

Here you will find daily **news and tutorials about R**, contributed by over 750 bloggers. There are many ways to **follow us -**

[By e-mail:](#)


Your e-mail here

 Subscribe

 45931 readers

BY FEEDBURNER

[On Facebook:](#)



If you are an R blogger yourself you are invited to [add your own R content feed to this site](#) (Non-English R bloggers should add themselves- [here](#))

[Jobs for R-users](#)

- [Financial Analyst/Modeler @ Mesa, Arizona, U.S.](#)
- [Research volunteer in Cardiac Surgery @ Philadelphia, Pennsylvania, U.S.](#)
- [Fit GLM distribution models to data using R](#)
- [Financial services startup looking for freelance Shiny/R/UI developer](#)
- [Data Scientists – PhD Paradise –](#)



[Germany](#)

Popular Searches

- [eclipse](#)

Recent Posts

- [Tutorial: Publish an R function as a SQL Server stored procedure with the sqldrutils package](#)
- [Soccer data sparring: Scraping, merging and analyzing exercises](#)
- [Seeking guidance in choosing and evaluating R packages](#)
- [Pets or livestock? Naming your RMarkdown chunks](#)
- [How to make best use of the byte compiler in R](#)
- [Bayesian Simple Linear Regression with Gibbs Sampling in R](#)
- [US Immigration Enforcement – Part 1](#)
- [Determine the CV of a Calculated Lab Reportable – Bioavailable Testosterone](#)
- [Exploring Assumptions of K-means Clustering using R](#)
- [Implementing Parallel Processing in R](#)
- [Radial kernel Support Vector Classifier](#)
- [Elegant correlation table using xtable R package](#)
- [R and Interactive Graphics](#)
- [Geospatial Queries using Pymongo in R](#)
- [More on “The Part-Time R-User”](#)

Other sites

- [SAS blogs](#)
- [Jobs for R-users](#)

Useful dplyr Functions

[f](#)[t](#)[in](#)[G+](#)

(w/examples)

July 10, 2017

By [S. Richter-Walsh](#)

 Like 139  Share  Share 12

(This article was first published on [Environmental Science and Data Analytics](#), and kindly contributed to [R-bloggers](#))

135
SHARES

 Share

 Tweet

The R package *dplyr* is an extremely useful resource for data cleaning, manipulation, visualisation and analysis. It contains a large number of very useful functions and is, without doubt, one of my top 3 R packages today (*ggplot2* and *reshape2* being the others). When I was learning how to use dplyr for the first time, I used [DataCamp](#) which offers some fantastic interactive courses on R. I also made use of a decent data wrangling cheat sheet which can be found [here](#).

There are many useful functions contained within the dplyr package. This post does not attempt to cover them all but does look at the major functions that are commonly used in data manipulation tasks. These are:

```
select()
filter()
mutate()
group_by()
summarise()
arrange()
join()
```

The data used in this post are taken from the UCI Machine Learning Repository and contain census information from 1994 for the USA. The dataset can be used for classification of income class in a machine learning setting and can be obtained [here](#).

```
require(dplyr)

# Data file
file <- "https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data"

# Some sensible variable names
df_names <- c("age", "wrkclass", "fnlweight", "education_lvl", "edu_score",
             "marital_status", "occupation", "relationship", "ethnic", "gender",
             "cap_gain", "cap_loss", "hrs_wk", "nationality", "income")

# Import the data
df <- read.csv(file, header = F,
              sep = ",",
              na.strings = c(" ", " ", " "),
              row.names = NULL,
              col.names = df_names)
```

Many data manipulation tasks in dplyr can be performed with the assistance of the forward-pipe operator (%>%). The pipe was first introduced in the *magrittr* package and has since been included in the dplyr package. It is an incredibly useful tool for fluid data manipulation and results in highly readable code.

The census dataset requires a bit of preprocessing to get it ready for classification algorithms. This post does not cover preprocessing nor does it include predictive modelling.

The first function I would like to introduce removes duplicate entries which, in fact, is a preprocessing step one may carry out in a data analysis. It is so useful that it must be included.

```
# Remove duplicate rows and check number of rows
df %>% distinct() %>% nrow()

# Drop duplicate rows and assign to new dataframe object
df_clean <- df %>% distinct()

# Drop duplicates based on one or more variables
df %>% distinct(gender, .keep_all = T)
df %>% distinct(gender, education_lvl, .keep_all = T)
```

Taking random samples of data is easy with dplyr.

```
# Sample random rows with or without replacement
sample_n(df, size = nrow(df) * 0.7, replace = F)
sample_n(df, size = 20, replace = T)

# Sample a proportion of rows with or without replacement
```

f
t
in
G+

```
sample_frac(df, size = 0.7, replace = F)
sample_frac(df, size = 0.8, replace = T)
```

Renaming variables is also easy with dplyr.

```
# Rename one or more variables in a dataframe
df <- df %>%
  rename("INCOME" = "income")

df <- df %>%
  rename("INCOME" = "income", "AGE" = "age")
```

The main “verbs” of dplyr are now introduced. Let’s begin with the **select()** verb which filters a dataframe by column.

```
# Select specific columns (note that INCOME is the new name from earlier)
df %>%
  select(education_lvl, INCOME)

# With dplyr 0.7.0 the pull() function extracts a variable as a vector
df %>%
  pull(age)

# Drop a column using the - operator (variable can be referenced by name or column position)
df %>%
  select(-edu_score)

df %>%
  select(-1, -4)

df %>%
  select(-c(2:6))
```

Some useful helper functions are available in dplyr and can be used in conjunction with the select() verb. Here are some quick examples.

```
# Select columns with their names starting with "e"
df %>%
  select(starts_with("e"))

# The negative sign works for dropping here too
df %>%
  select(-starts_with("e"))

# Select columns with some pattern in the column name
df %>%
  select(contains("edu"))

# Reorder data to place a particular column at the start followed by all others using everything()
df %>%
  select(INCOME, everything())

# Select columns ending with a pattern
df %>%
  select(ends_with("e"))

df %>%
  select(ends_with("_loss"))
```

The next major verb we look at is **filter()** which, surprisingly enough, filters a dataframe by row based on one or more conditions.

```
# Filter rows to retain observations where age is greater than 30
df %>%
  filter(age > 30)

# Filter by multiple conditions using the %in% operator (make sure strings match)
df %>%
  filter(relationship %in% c(" Unmarried", " Wife"))

# You can also use the OR operator (|)
df %>%
  filter(relationship == " Husband" | relationship == " Wife")

# Filter using the AND operator
df %>%
  filter(age > 30 & INCOME == " >50K")

# Combine them too
df %>%
  filter(education_lvl %in% c(" Doctorate", " Masters") & age > 30)

# The NOT condition (filter out doctorate holders)
df %>%
  filter(education_lvl != " Doctorate")

# The grepl() function can be conveniently used with filter()
df %>%
  filter(grepl(" Wi", relationship))
```

Next, we look at the **summarise()** verb which allows one to dynamically summarise groups of data and even pipe groups to ggplot data visualisations.

```
# The summarise() verb in dplyr is useful for summarising grouped data
df %>%
```



```

filter(INCOME == " >50K") %>%
summarise(mean_age = mean(age),
           median_age = median(age),
           sd_age = sd(age))

# Summarise multiple variables using summarise_at()
df %>%
  filter(INCOME == " >50K") %>%
  summarise_at(vars(age, hrs_wk),
               funs(n(),
                    mean,
                    median))

# We can also summarise with custom functions
# The . in parentheses represents all called variables
df %>%
  summarise_at(vars(age, hrs_wk),
               funs(n(),
                    missing = sum(is.na(.)),
                    mean = mean(., na.rm = T)))

# Create a new summary statistic with an anonymous function
df %>%
  summarise_at(vars(age),
               function(x) { sum((x - mean(x)) / sd(x)) })

# Summarise conditionally using summarise_if()
df %>%
  filter(INCOME == " >50K") %>%
  summarise_if(is.numeric,
               funs(n(),
                    mean,
                    median))

# Subset numeric variables and use summarise_all() to get summary statistics
ints <- df[sapply(df, is.numeric)]
summarise_all(ints,
               funs(mean,
                    median,
                    sd,
                    var))

```

Next up is the **arrange()** verb which is useful for sorting data in ascending or descending order (ascending is default).

```

# Sort by ascending age and print top 10
df %>%
  arrange(age) %>%
  head(10)

# Sort by descending age and print top 10
df %>%
  arrange(desc(age)) %>%
  head(10)

```

The **group_by()** verb is useful for grouping together observations which share common characteristics.

```

# The group_by verb is extremely useful for data analysis
df %>%
  group_by(gender) %>%
  summarise(Mean = mean(age))

df %>%
  group_by(relationship) %>%
  summarise(total = n())

df %>%
  group_by(relationship) %>%
  summarise(total = n(),
            mean_age = mean(age))

```

The **mutate()** verb is used to create new variables from existing local variables or global objects. New variables, such as sequences, can be also specified within mutate().

```

# Create new variables from existing or global variables
df %>%
  mutate(norm_age = (age - mean(age)) / sd(age))

# Multiply each numeric element by 1000 (the name "new" is added to the original variable name)
df %>%
  mutate_if(is.numeric,
            funs(new = (. * 1000))) %>%
  head()

```

The **join()** verb is used to merge rows from disjoint tables which share a primary key ID or some other common variable. There are many join variants but I will consider just left, right, inner and full joins.

```

# Create ID variable which will be used as the primary key
df <- df %>%
  mutate(ID = seq(1:nrow(df))) %>%
  select(ID, everything())

```



```
# Create two tables (purposely overlap to facilitate joins)
table_1 <- df[1:50 , ] %>%
  select(ID, age, education_lvl)

table_2 <- df[26:75 , ] %>%
  select(ID, gender, INCOME)

# Left join joins rows from table 2 to table 1 (the direction is implicit in the argument order)
left_join(table_1, table_2, by = "ID")

# Right join joins rows from table 1 to table 2
right_join(table_1, table_2, by = "ID")

# Inner join joins and retains only complete cases
inner_join(table_1, table_2, by = "ID")

# Full join joins and retains all values
full_join(table_1, table_2, by = "ID")
```

That wraps up a brief demonstration of some of dplyr's excellent functions. For additional information on the functions and their arguments, check out the help documentation using the template: ?

References

Hadley Wickham, Romain Francois, Lionel Henry and Kirill Müller (2017). dplyr: A Grammar of Data Manipulation. R package version 0.7.0.
<https://CRAN.R-project.org/package=dplyr>

H. Wickham. ggplot2: Elegant Graphics for Data Analysis. Springer-Verlag New York, 2009.



Related

[Data Manipulation in R with dplyr – Part 1](#)
 dplyr is one of the packages in R that makes R so loved by data scientists. It has three main goals:
 December 17, 2015
 In "R bloggers"

[Enernoc smart meter data - forecast electricity consumption with similar day approach in R](#)
 Deployment of smart grids gives space to an occurrence of new methods of machine learning and data
 November 11, 2016
 In "R bloggers"



[Tutorials for learning R](#)
 December 10, 2015
 In "R guest posts"

135
SHARES

Share

Tweet

To leave a comment for the author, please follow the link and comment on their blog:
[Environmental Science and Data Analytics](#).

R-bloggers.com offers [daily e-mail updates](#) about R news and [tutorials](#) on topics such as: [Data science](#), [Big Data](#), [R jobs](#), visualization ([ggplot2](#), [Boxplots](#), [maps](#), [animation](#)), programming ([RStudio](#), [Sweave](#), [LaTeX](#), [SQL](#), [Eclipse](#), [git](#), [hadoop](#), [Web Scraping](#)) statistics ([regression](#), [PCA](#), [time series](#), [trading](#)) and more...

If you got this far, why not **subscribe for updates** from the site?
 Choose your flavor: [e-mail](#), [twitter](#), [RSS](#), or [facebook](#)...

Like 139 Share

Share 12

Comments are closed.

Search R-bloggers

Custom Search

Recent popular posts

- [Elegant correlation table using xtable R](#)

[package](#)

- [Exploring Assumptions of K-means Clustering using R](#)
- [R and Interactive Graphics](#)
- [Bayesian Simple Linear Regression with Gibbs Sampling in R](#)

Most visited articles of the week


1. [How to write the first for loop in R](#)
2. [Installing R packages](#)
3. [Using apply, sapply, lapply in R](#)
4. [Tutorials for learning R](#)
5. [How to perform a Logistic Regression in R](#)
6. [In-depth introduction to machine learning in 15 hours of expert videos](#)
7. [Elegant correlation table using xtable R package](#)
8. [Exploring Assumptions of K-means Clustering using R](#)
9. [How to Make a Histogram with Basic R](#)

Sponsors



Want to find out
how to use R in a
regulated industry?

www.mango-solutions.com



 DataCamp


Sign Up Now

RStudio Connect

Push-Button Publishing | Scheduled Updates
For all the work you create in R



Free Evaluation



Highland Statistics Ltd.

Course:

Introduction to Regression
Models with Spatial &
Temporal Correlation

Southampton, October 2017
Banff, December 2017

Register for free updates
in all areas of statistics



STATISTICS
VIEWS



Bringing Statistics Together

WILEY


f
t
in
G+




Quantide: statistical consulting and training



START



f
t
in
G+



[Contact us](#) if you wish to help support R-bloggers, and place **your banner here**.

[Jobs for R users](#)

- [Financial Analyst/Modeler @ Mesa, Arizona, U.S.](#)
- [Research volunteer in Cardiac Surgery @ Philadelphia, Pennsylvania, U.S.](#)
- [Fit GLM distribution models to data using R](#)
- [Financial services startup looking for freelance Shiny/R/UI developer](#)
- [Data Scientists – PhD Paradise – Germany](#)
- [Technical Project Manager](#)
- [Software Developer](#)

[Full list of contributing R-bloggers](#)

[R-bloggers](#) was founded by [Tal Galili](#), with gratitude to the [R](#) community.

Is powered by [WordPress](#) using a [bavotasan.com](#) design.

Copyright © 2017 **R-bloggers**. All Rights Reserved. [Terms and Conditions](#) for this website

f
t
in
G+